

Black_Friday_V2

March 27, 2023

Read & Import packages.

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import mean_squared_error, mean_absolute_percentage_error, \
    r2_score

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost.sklearn import XGBRegressor
```

Master Dataset contains total of 7,83,667 records.

We divided the data in the ratio of 70-30%

```
[5]: df_train = pd.read_csv('https://raw.githubusercontent.com/commonemail270/EDA/
    ↪main/Balck_friday_dataset/train.csv')
df_test = pd.read_csv('https://raw.githubusercontent.com/commonemail270/EDA/
    ↪main/Balck_friday_dataset/test.csv')
df_train.head()
```

```
[5]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
0	1000001	P00069042	F	0-17	10	A	
1	1000001	P00248942	F	0-17	10	A	
2	1000001	P00087842	F	0-17	10	A	
3	1000001	P00085442	F	0-17	10	A	
4	1000002	P00285442	M	55+	16	C	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
0	2	0	3	
1	2	0	1	
2	2	0	12	
3	2	0	12	

4		4+	0	8
	Product_Category_2	Product_Category_3	Purchase	
0	NaN	NaN	8370	
1	6.0	14.0	15200	
2	NaN	NaN	1422	
3	14.0	NaN	1057	
4	NaN	NaN	7969	

```
[6]: df_test.head()
```

```
[6]:   User_ID Product_ID Gender   Age Occupation City_Category \
0  1000004  P00128942     M  46-50           7           B
1  1000009  P00113442     M  26-35          17           C
2  1000010  P00288442     F  36-45           1           B
3  1000010  P00145342     F  36-45           1           B
4  1000011  P00053842     F  26-35           1           C

   Stay_In_Current_City_Years  Marital_Status  Product_Category_1 \
0                             2                1                  1
1                             0                0                  3
2                             4+               1                  5
3                             4+               1                  4
4                             1                0                  4

   Product_Category_2  Product_Category_3
0                   11.0                 NaN
1                   5.0                 NaN
2                   14.0                 NaN
3                   9.0                 NaN
4                   5.0                 12.0
```

Data Analysis

```
[7]: df_train.isnull().sum().sort_values(ascending=False) * 100 / len(df_train)
```

```
[7]: Product_Category_3      69.672659
Product_Category_2      31.566643
User_ID                  0.000000
Product_ID               0.000000
Gender                   0.000000
Age                     0.000000
Occupation               0.000000
City_Category            0.000000
Stay_In_Current_City_Years  0.000000
Marital_Status           0.000000
Product_Category_1       0.000000
```

```
Purchase                                0.000000
dtype: float64
```

```
[170]: df_test.isnull().sum().sort_values(ascending=False) * 100 / len(df_test)
```

```
[170]: Product_Category_3      69.590195
      Product_Category_2      30.969311
      Product_Category_1       0.000000
      Marital_Status          0.000000
      Stay_In_Current_City_Years  0.000000
      City_Category           0.000000
      Occupation              0.000000
      Age                     0.000000
      Gender                  0.000000
      Product_ID              0.000000
      User_ID                 0.000000
dtype: float64
```

```
[10]:
```

```
[10]: 233599
```

```
[171]: df_train.describe().T
```

```
[171]:
```

	count	mean	std	min	25%	\
User_ID	550068.0	1.003029e+06	1727.591586	1000001.0	1001516.0	
Occupation	550068.0	8.076707e+00	6.522660	0.0	2.0	
Marital_Status	550068.0	4.096530e-01	0.491770	0.0	0.0	
Product_Category_1	550068.0	5.404270e+00	3.936211	1.0	1.0	
Product_Category_2	376430.0	9.842329e+00	5.086590	2.0	5.0	
Product_Category_3	166821.0	1.266824e+01	4.125338	3.0	9.0	
Purchase	550068.0	9.263969e+03	5023.065394	12.0	5823.0	

	50%	75%	max
User_ID	1003077.0	1004478.0	1006040.0
Occupation	7.0	14.0	20.0
Marital_Status	0.0	1.0	1.0
Product_Category_1	5.0	8.0	20.0
Product_Category_2	9.0	15.0	18.0
Product_Category_3	14.0	16.0	18.0
Purchase	8047.0	12054.0	23961.0

We will impute the missing values of Product Category 2 with mean value as Mean and Median of feature is close to 9.

```
[172]: df_train['Product_Category_2'] = df_train['Product_Category_2'].
      ↪ fillna(df_train['Product_Category_2'].mean())
```

```
[173]: df_test['Product_Category_2'] = df_test['Product_Category_2'].
        ↪fillna(df_test['Product_Category_2'].mean())
```

```
[174]: df_train['Stay_In_Current_City_Years'].unique()
```

```
[174]: array(['2', '4+', '3', '1', '0'], dtype=object)
```

```
[175]: df_train['Age'].unique()
```

```
[175]: array(['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25'],
            dtype=object)
```

```
[176]: df_train['Gender'].unique()
```

```
[176]: array(['F', 'M'], dtype=object)
```

```
[177]: df_train['City_Category'].unique()
```

```
[177]: array(['A', 'C', 'B'], dtype=object)
```

****Data Cleaning.**

```
[178]: train_df = df_train.copy()
        train_df.Stay_In_Current_City_Years.replace({'0':0,
                                                    '1':1,
                                                    '2':2,
                                                    '3':3,
                                                    '4+':4},inplace=True)
        train_df.Age.replace({"55+": "55 & above"},inplace=True)
        train_df.Gender.replace({'M':1, 'F':0},inplace=True)
        train_df.City_Category.replace({'A':1, 'B':2, 'C':3},inplace=True)
```

```
[179]: train_df = train_df.rename(columns={'Age': 'Age_Groups'})
        train_df.head()
```

```
[179]:
```

	User_ID	Product_ID	Gender	Age_Groups	Occupation	City_Category	\
0	1000001	P00069042	0	0-17	10		1
1	1000001	P00248942	0	0-17	10		1
2	1000001	P00087842	0	0-17	10		1
3	1000001	P00085442	0	0-17	10		1
4	1000002	P00285442	1	55 & above	16		3

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
0	2	0		3
1	2	0		1
2	2	0		12
3	2	0		12

	4	4	0	8
	Product_Category_2	Product_Category_3	Purchase	
0	9.842329	NaN	8370	
1	6.000000	14.0	15200	
2	9.842329	NaN	1422	
3	14.000000	NaN	1057	
4	9.842329	NaN	7969	

```
[180]: test_df = df_test.copy()
test_df.Stay_In_Current_City_Years.replace({'0':0,
                                             '1':1,
                                             '2':2,
                                             '3':3,
                                             '4+':4},inplace=True)
test_df.Age.replace({"55+": "55 & above"},inplace=True)
test_df.Gender.replace({'M':1, 'F':0},inplace=True)
test_df.City_Category.replace({'A':1, 'B':2, 'C':3},inplace=True)
```

```
[181]: test_df = test_df.rename(columns={'Age': 'Age_Groups'})
test_df.head(1)
```

```
[181]: User_ID Product_ID Gender Age_Groups Occupation City_Category \
0 1000004 P00128942 1 46-50 7 2

Stay_In_Current_City_Years Marital_Status Product_Category_1 \
0 2 1 1

Product_Category_2 Product_Category_3
0 11.0 NaN
```

****Exploratory Data Analysis.**

```
[182]: df_train.head(2)
```


```
[182]: User_ID Product_ID Gender Age Occupation City_Category \
0 1000001 P00069042 F 0-17 10 A
1 1000001 P00248942 F 0-17 10 A

Stay_In_Current_City_Years Marital_Status Product_Category_1 \
0 2 0 3
1 2 0 1

Product_Category_2 Product_Category_3 Purchase
0 9.842329 NaN 8370
1 6.000000 14.0 15200
```

Univariate Analysis.

```
[183]: plt.title("Gender Distribution")
sns.countplot(df_train['Gender'], palette='viridis', order = df_train['Gender'].
    ↪ value_counts().index)
plt.show()
```




Insights:

1. The Gender feature has high data imbalance. The ratio of count of female customers is very less compared to male customers.
2. Need to handle this class imbalance using SMOTE/Oversampling techniques.

```
[184]: plt.figure(figsize = (10,6))
plt.title("City Category Distribution")
palette_color = sns.color_palette('coolwarm')

# plotting data on chart
plt.pie(df_train['City_Category'].value_counts(normalize=True),
    ↪ colors=palette_color, autopct='%0f%%', labels =
    ↪ list(df_train['City_Category'].value_counts().index))


# displaying chart
plt.show()
```



Insight:

1. City category B has observed to be having highest percentage of customers who have purchased in Black Friday Sale i.e 42% compared to A and C.

```
[185]: plt.title("Marital Status Distribution")
sns.countplot(df_train['Marital_Status'], palette='viridis')
plt.show()
```



Insight:

1. Data shows unmarried customers have spent on Black Friday sale more than married customers.

```
[186]: plt.figure(figsize = (10,6))
plt.title("Age Group Wise Distribution")


my_circle = plt.Circle((0, 0), 0.7, color='white')
palette_color = sns.color_palette('coolwarm')

# plotting data on chart
plt.pie(df_train['Age'].value_counts(normalize=True), colors=palette_color,
        autopct='%.0f%%', labels = list(df_train['Age'].value_counts().index))

p = plt.gcf()
p.gca().add_artist(my_circle)

# displaying chart
plt.show()
```


Age Group Wise Distribution




Insights:

1. Age group 26-35 years has highest percentage of customers who have spent on Black Friday sale i.e almost around 40%.
2. Age group 0-17 years has lowest percentage of customers who have spent (3%). And it is reasonable as teenage customers are less probable to have income.
3. Age group 18-25 & 36-45 years has average percentage of customers who spent on sale (around 20%).
4. Customers with Age above 45 years has observed to have decreasing percentage of customers as trend.

```
[187]: plt.figure(figsize = (10,6))  
plt.title("Purchase Distribution")  
sns.distplot(df_train['Purchase'])
```

```
[187]: <matplotlib.axes._subplots.AxesSubplot at 0x17656552c10>
```




Insights:

1. The important observation from the above visualisation can be made that there are some outliers present in the dependent/target feature "Purchase".
2. According to distplot data is nearly normally distributed.

Bivariate Analysis.

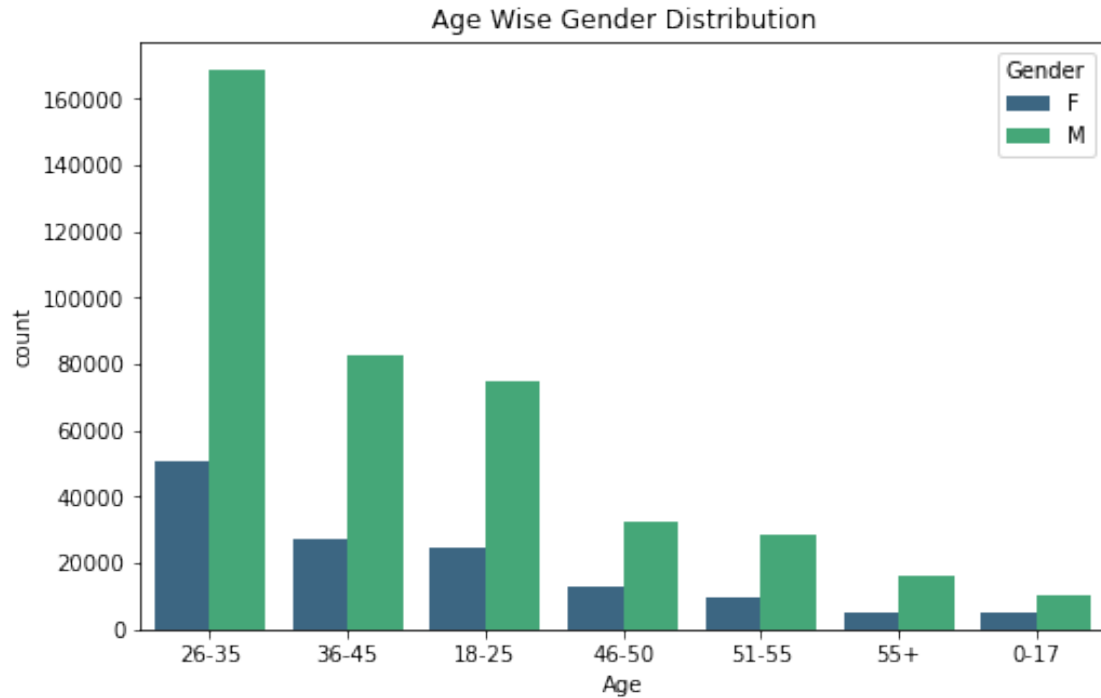
```
[188]: plt.figure(figsize = (8,5))
plt.title("Age Wise Marital Status Distribution")
sns.countplot(df_train['Age'], hue=df_train['Marital_Status'],
              palette='viridis', order = df_train['Age'].value_counts().index)
plt.show()
```



Insights:

1. Age group 0-17 years has all the single customers.
2. Age group 18-25 & 36-45 years has high single ratio than married.
3. Age group 26-35 years has highest ratio of both being single and married customers.
4. As age group is getting increased the ratio of being single is reduced. For example, 46-50, 51-55 & 55+, etc.

```
[189]: plt.figure(figsize = (8,5))
plt.title("Age Wise Gender Distribution")
sns.countplot(df_train['Age'], hue=df_train['Gender'], palette='viridis', order_
↳ df_train['Age'].value_counts().index)
plt.show()
```




Insights:

1. In each Age group Male customers are dominating in spending/purchase in Black Friday Sale.
2. Age group 26-35 years has highest number of customers, whereas group 18-25 & 36-45 years has average number of customers.
3. Age groups 0-17 & 55+ years has lowest number of customers.
4. Less number of customers are witnessed in age groups 46-50 & 51-55 years to purchase in Sale.

```
[190]: plt.figure(figsize = (12,6))
plt.title("Occupation Wise Purchase Distribution")
sns.barplot(x='Occupation', y='Purchase', data=df_train)
```

```
[190]: <matplotlib.axes._subplots.AxesSubplot at 0x17656e6ce50>
```




Insight:

1. Almost all of the Occupation categories have spent 8000-10000 in Black Friday sale

```
[191]: plt.figure(figsize = (12,6))
plt.title("Product Category 1 Wise Purchase Distribution")
sns.barplot(x='Product_Category_1', y='Purchase', data=df_train)
```

[191]: <matplotlib.axes._subplots.AxesSubplot at 0x1765ecfe1c0>




Insights:

1. Product category 10 has highest Purchase happened in Black Friday Sale.
2. Other Categories has a dispersed sales.

```
[192]: plt.figure(figsize = (15,6))
plt.title("Product Category 2 Wise Purchase Distribution")
sns.barplot(x='Product_Category_2', y='Purchase', data=df_train)
```

```
[192]: <matplotlib.axes._subplots.AxesSubplot at 0x17663548d00>
```



Insights:

1. Product category 10 has highest Purchase happened in Black Friday Sale.
2. Other Categories has a dispersed sales.
3. We can't make any solid statement from above visualizations.

Multivariate Analysis.

```
[193]: pd.crosstab([df_train['Gender']], [df_train['Marital_Status'],
df_train['Age']], normalize=True)*100
```

```
[193]: Marital_Status      0
Age
Gender
F          0.924068    3.337224    5.468051    3.026717    0.575565    0.650829
M          1.821411   10.941738   18.764589    9.040337    1.731422    1.319655

Marital_Status      1
Age
55+    18-25    26-35    36-45    46-50    51-55
```

Gender						
F	0.346866	1.140041	3.758444	1.912673	1.823956	1.147858
M	1.086229	2.698757	11.928889	6.020165	4.177302	3.880975

Marital_Status	
Age	55+
Gender	
F	0.577201
M	1.899038

Insights:

1. In all unmarried customers, Age groups 18-25 & 26-35 years has highest percentage of customers in both genders Male & Female.
2. In all Male married customers, Age groups 18-25 & 26-35 years has highest percentage of customers.
3. In all Female married customers, Age groups 26-35 & 36-45 years has highest purchase customers.

```
[194]: pd.crosstab([df_train['City_Category'],df_train['Marital_Status']],
    ↪[df_train['Gender'], df_train['Age']], normalize=True)*100
```

```
[194]: Gender
Age
City_Category Marital_Status
A
0 0.263058 0.861348 1.816866 0.853894
1 0.000000 0.278329 1.362922 0.437764
B
0 0.284510 1.557989 2.550594 1.209305
1 0.000000 0.566475 1.346561 0.810445
C
0 0.376499 0.917887 1.100591 0.963517
1 0.000000 0.295236 1.048961 0.664463
```

Gender					M \
Age		46-50	51-55	55+	0-17
City_Category	Marital_Status				
A	0	0.048358	0.085989	0.038541	0.199430
	1	0.178887	0.237243	0.027633	0.000000
B	0	0.269421	0.327778	0.077263	0.703549
	1	0.894798	0.443582	0.168343	0.000000
C	0	0.257786	0.237062	0.231062	0.918432
	1	0.750271	0.467033	0.381226	0.000000

Gender					\
Age		18-25	26-35	36-45	46-50
City_Category	Marital_Status				
A	0	3.100526	6.613728	1.822684	0.394497
	1	0.765542	3.613008	1.724514	0.761179

B	0	4.561800	7.664689	4.268200	0.612833
	1	1.175855	5.087735	2.365162	1.932670
C	0	3.279413	4.486173	2.949454	0.724092
	1	0.757361	3.228146	1.930489	1.483453


Gender			
Age		51-55	55+
City_Category	Marital_Status		
A	0	0.237243	0.238698
	1	0.548296	0.344685
B	0	0.539751	0.227608
	1	1.914127	0.465215
C	0	0.542660	0.619923
	1	1.418552	1.089138

Insights:

1. City B has highest percentage of Female customers in Age group 26-35 years in both Married and Unmarried marital status.
2. City B has highest percentage of Male customers in Age group 26-35 years in both Married and Unmarried marital status.
3. The above tabular representation shows that, in all the city categories A, B & C, the customers who have purchased/spent on Black Friday Sale always have high percentage of Unmarried customers irrespective of their Gender, Age groups.
4. Hence, we can say that Unmarried customers are more tend to spend in Black Friday sale.

```
[195]: plt.figure(figsize = (12,6))
plt.title("Occupation Wise Purchase Distribution")
sns.barplot(x='Occupation', y='Purchase', hue = 'Gender',data=df_train)
```

```
[195]: <matplotlib.axes._subplots.AxesSubplot at 0x1765f1214c0>
```





Insights:

1. In the Occupation level 18, Female customers has high purchase amount compared to Male customers.
2. The purchase amount of all the Occupation levels is almost in the range of 8000-10000, which indicates Occupation levels has no greater impact on the purchase in Black Friday..

```
[196]: plt.figure(figsize = (7,4))
plt.title("Correlation Matrix")
feats = ['Gender', 'City_Category', 'Occupation', 'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category_1', 'Product_Category_2', 'Product_Category_3', 'Purchase']
sns.heatmap(df_train[feats].corr(), annot = True)
```

```
[196]: <matplotlib.axes._subplots.AxesSubplot at 0x17656e03f10>
```



Insights:

1. Occupation levels has no impact/correlation on independent as well as dependent features. Occupation feature can be dropped based on the various performance comparisons of the Model.
2. Product Category 3 has no correlation with target feature (Purchase). Also it has 70% of missing values, Hence, Product Category 3 feature can be dropped
3. Product Category 1 & Product Category 2 has strong negative correlation with target feature (Purchase). Which means the change of value in one feature varies with change of value in other feature. This is called as Inverse correlation. In other words, If the value of Purchase Category 1 increases, that will result in reduction in Purchase Value.

****Feature Engineering & Extraction**

Transforming categorical feature into numeric using dummy encoding technique (One hot encoding).

```
[197]: one_hot = pd.
        ↳ get_dummies(data=train_df['Age_Groups'], prefix='Age_Group', drop_first=True)
        one_hot
```

```
[197]:
```

	Age_Group_18-25	Age_Group_26-35	Age_Group_36-45	Age_Group_46-50	\
0	0	0	0	0	
1	0	0	0	0	

2	0	0	0	0
3	0	0	0	0
4	0	0	0	0
...
550063	0	0	0	0
550064	0	1	0	0
550065	0	1	0	0
550066	0	0	0	0
550067	0	0	0	1

	Age_Group_51-55	Age_Group_55 & above
0	0	0
1	0	0
2	0	0
3	0	0
4	0	1
...
550063	1	0
550064	0	0
550065	0	0
550066	0	1
550067	0	0

[550068 rows x 6 columns]

```
[198]: train_df = pd.concat([train_df,one_hot],axis = 1).drop('Age_Groups',axis = 1)
train_df
```

```
[198]:
```

	User_ID	Product_ID	Gender	Occupation	City_Category	\
0	1000001	P00069042	0	10	1	
1	1000001	P00248942	0	10	1	
2	1000001	P00087842	0	10	1	
3	1000001	P00085442	0	10	1	
4	1000002	P00285442	1	16	3	
...	
550063	1006033	P00372445	1	13	2	
550064	1006035	P00375436	0	1	3	
550065	1006036	P00375436	0	15	2	
550066	1006038	P00375436	0	1	3	
550067	1006039	P00371644	0	0	2	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
0	2	0	3	
1	2	0	1	
2	2	0	12	
3	2	0	12	
4	4	0	8	

...
550063	1	1	20
550064	3	0	20
550065	4	1	20
550066	2	0	20
550067	4	1	20

	Product_Category_2	Product_Category_3	Purchase	Age_Group_18-25	\
0	9.842329	NaN	8370	0	
1	6.000000	14.0	15200	0	
2	9.842329	NaN	1422	0	
3	14.000000	NaN	1057	0	
4	9.842329	NaN	7969	0	
...	
550063	9.842329	NaN	368	0	
550064	9.842329	NaN	371	0	
550065	9.842329	NaN	137	0	
550066	9.842329	NaN	365	0	
550067	9.842329	NaN	490	0	

	Age_Group_26-35	Age_Group_36-45	Age_Group_46-50	Age_Group_51-55	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
...	
550063	0	0	0	1	
550064	1	0	0	0	
550065	1	0	0	0	
550066	0	0	0	0	
550067	0	0	1	0	

	Age_Group_55 & above
0	0
1	0
2	0
3	0
4	1
...	...
550063	0
550064	0
550065	0
550066	1
550067	0

[550068 rows x 17 columns]

```
[199]: one_hot_test = pd.
        get_dummies(data=test_df['Age_Groups'],prefix='Age_Group',drop_first=True)
one_hot_test
```

```
[199]:
```

	Age_Group_18-25	Age_Group_26-35	Age_Group_36-45	Age_Group_46-50	\
0	0	0	0	1	
1	0	1	0	0	
2	0	0	1	0	
3	0	0	1	0	
4	0	1	0	0	
...	
233594	0	1	0	0	
233595	0	1	0	0	
233596	0	1	0	0	
233597	0	0	0	1	
233598	0	0	0	1	

	Age_Group_51-55	Age_Group_55 & above
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...
233594	0	0
233595	0	0
233596	0	0
233597	0	0
233598	0	0

[233599 rows x 6 columns]

```
[200]: test_df = pd.concat([test_df,one_hot_test],axis = 1).drop('Age_Groups',axis = 1)
test_df.head()
```

```
[200]:
```

	User_ID	Product_ID	Gender	Occupation	City_Category	\
0	1000004	P00128942	1	7	2	
1	1000009	P00113442	1	17	3	
2	1000010	P00288442	0	1	2	
3	1000010	P00145342	0	1	2	
4	1000011	P00053842	0	1	3	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
0	2	1	1	
1	0	0	3	
2	4	1	5	
3	4	1	4	

	4	1	0	4
	Product_Category_2	Product_Category_3	Age_Group_18-25	Age_Group_26-35 \
0	11.0	NaN	0	0
1	5.0	NaN	0	1
2	14.0	NaN	0	0
3	9.0	NaN	0	0
4	5.0	12.0	0	1

	Age_Group_36-45	Age_Group_46-50	Age_Group_51-55	Age_Group_55 & above
0	0	1	0	0
1	0	0	0	0
2	1	0	0	0
3	1	0	0	0
4	0	0	0	0

Drop unwanted features.

Dropping Product_Category_3 as it has 70% of the data missing..

```
[201]: train_df.drop(columns = ['User_ID', 'Product_ID', 'Product_Category_3',
    ↳ 'Occupation'], axis = 1, inplace=True)
train_df
```

```
[201]:
```

	Gender	City_Category	Stay_In_Current_City_Years	Marital_Status \
0	0	1	2	0
1	0	1	2	0
2	0	1	2	0
3	0	1	2	0
4	1	3	4	0
...
550063	1	2	1	1
550064	0	3	3	0
550065	0	2	4	1
550066	0	3	2	0
550067	0	2	4	1

	Product_Category_1	Product_Category_2	Purchase	Age_Group_18-25 \
0	3	9.842329	8370	0
1	1	6.000000	15200	0
2	12	9.842329	1422	0
3	12	14.000000	1057	0
4	8	9.842329	7969	0
...
550063	20	9.842329	368	0
550064	20	9.842329	371	0
550065	20	9.842329	137	0
550066	20	9.842329	365	0

550067		20	9.842329	490	0
	Age_Group_26-35	Age_Group_36-45	Age_Group_46-50	Age_Group_51-55	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
...	
550063	0	0	0	1	
550064	1	0	0	0	
550065	1	0	0	0	
550066	0	0	0	0	
550067	0	0	1	0	

	Age_Group_55 & above
0	0
1	0
2	0
3	0
4	1
...	...
550063	0
550064	0
550065	0
550066	1
550067	0

[550068 rows x 13 columns]

```
[202]: test_df.drop(columns = ['User_ID', 'Product_ID', 'Product_Category_3', 'Occupation'],axis = 1,inplace=True)
test_df
```

```
[202]:
```

	Gender	City_Category	Stay_In_Current_City_Years	Marital_Status	\
0	1	2	2	1	
1	1	3	0	0	
2	0	2	4	1	
3	0	2	4	1	
4	0	3	1	0	
...	
233594	0	2	4	1	
233595	0	2	4	1	
233596	0	2	4	1	
233597	0	3	4	0	
233598	0	2	4	1	

	Product_Category_1	Product_Category_2	Age_Group_18-25	\
0	1	11.000000	0	
1	3	5.000000	0	
2	5	14.000000	0	
3	4	9.000000	0	
4	4	5.000000	0	
...	
233594	8	9.849586	0	
233595	5	8.000000	0	
233596	1	5.000000	0	
233597	10	16.000000	0	
233598	4	5.000000	0	

	Age_Group_26-35	Age_Group_36-45	Age_Group_46-50	Age_Group_51-55	\
0	0	0	1	0	
1	1	0	0	0	
2	0	1	0	0	
3	0	1	0	0	
4	1	0	0	0	
...	
233594	1	0	0	0	
233595	1	0	0	0	
233596	1	0	0	0	
233597	0	0	1	0	
233598	0	0	1	0	

	Age_Group_55 & above
0	0
1	0
2	0
3	0
4	0
...	...
233594	0
233595	0
233596	0
233597	0
233598	0

[233599 rows x 12 columns]

**Outlier treatment on train dataset.

```
[203]: import scipy.stats as stats

train_df['zscore'] = stats.zscore(train_df['Purchase'])
```



```
print("The Minimum zscore of feature Purchase is:", np.round(train_df['zscore'].
    ↪min()))
print("The Maximum zscore of feature Purchase is:", np.round(train_df['zscore'].
    ↪max()))
train_df.drop('zscore',axis = 1, inplace = True)
```

The Minimum zscore of feature Purchase is: -2.0

The Maximum zscore of feature Purchase is: 3.0

As per Z_score method to detect and remove outliers, the threshold value of lambda(the decision parameter) should be +3 & -3. And as our minimum and maximum z_score is below the threshold, we can say that there are no outliers.

****Separate Independent and dependent features & Split the data in train and test sets.**

```
[205]: X = train_df.drop("Purchase",axis = True)
y = train_df["Purchase"]

x_train, x_test, y_train, y_test=train_test_split(X,y, test_size=0.2,
    ↪random_state=101)
```

****Modeling.**

Linear Regression.

```
[206]: lr = LinearRegression()
lr.fit(x_train, y_train)

y_pred_lr = lr.predict(x_test)

print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred_lr)))
print('MAPE:', metrics.mean_absolute_percentage_error(y_test, y_pred_lr))
print('R2 Score:',metrics.r2_score(y_test,y_pred_lr)*100)
```

RMSE: 4694.569599682641

MAPE: 1.1134022216354589

R2 Score: 12.631382780679578

```
[207]: y_pred_train_lr = lr.predict(x_train)

print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred_train_lr)))
print('MAPE:', metrics.mean_absolute_percentage_error(y_train, y_pred_train_lr))
print('R2 Score:',metrics.r2_score(y_train,y_pred_train_lr)*100)
```

RMSE: 4690.693640109924

MAPE: 1.1247366725912407

R2 Score: 12.800269305429568

Random Forest Regressor.

```
[208]: rfr=RandomForestRegressor(random_state=0)
rfr.fit(x_train, y_train)

y_pred_rfr=rfr.predict(x_test)

print("RMSE:", np.sqrt(metrics.mean_squared_error(y_test, y_pred_rfr)))
print('MAPE:', metrics.mean_absolute_percentage_error(y_test, y_pred_rfr))
print("R2 Score:",round(metrics.r2_score(y_test,y_pred_rfr)*100,2))
```

RMSE: 3015.3225343046383
MAPE: 0.3408985831052717
R2 Score: 63.96

```
[209]: y_pred_train_rfr = rfr.predict(x_train)

print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred_train_rfr)))
print('MAPE:', metrics.mean_absolute_percentage_error(y_train,
↳y_pred_train_rfr))
print('R2 Score:',metrics.r2_score(y_train,y_pred_train_rfr)*100)
```

RMSE: 2853.0580891810814
MAPE: 0.32484695245946726
R2 Score: 67.74015817141503

XGBoost Regressor.

```
[210]: xgb_reg=XGBRegressor()
xgb_reg.fit(x_train,y_train)
y_pred_xgb = xgb_reg.predict(x_test)
print("RMSE:",np.sqrt(metrics.mean_squared_error(y_test, y_pred_xgb)))
print('MAPE:', metrics.mean_absolute_percentage_error(y_test, y_pred_xgb))
print("R2 Score:",round(metrics.r2_score(y_test,y_pred_xgb)*100,2))
```

RMSE: 2972.3568267550327
MAPE: 0.3567527385887654
R2 Score: 64.98

```
[211]: y_pred_train_xgb = xgb_reg.predict(x_train)

print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred_train_xgb)))
print('MAPE:', metrics.mean_absolute_percentage_error(y_train,
↳y_pred_train_xgb))
print('R2 Score:',metrics.r2_score(y_train,y_pred_train_xgb)*100)
```

RMSE: 2935.1009550901113
MAPE: 0.3530069512521715
R2 Score: 65.85814662280279

Performance Metrics Comparison:

RMSE: 4694.56 MAPE: 1.11 R2 Score: 12.63

RMSE: 3015.32 MAPE: 0.34 R2 Score: 63.96

RMSE: 2972.35 MAPE: 0.35 R2 Score: 64.98

Clearly, XGBoost Regressor is performing well as we are getting lowest RSME, low MAPE and high R-squared errors as compared to other models.

**Merge outputs with x_test and compare Actual and Predicted Values.

```
[212]: output_df = df_train[df_train.index.isin(x_test.index)]
output_df.head(2)
```

```
[212]:      User_ID Product_ID Gender   Age  Occupation City_Category \
2    1000001  P00087842      F  0-17          10           A
12   1000005  P00031342      M 26-35          20           A

      Stay_In_Current_City_Years  Marital_Status  Product_Category_1 \
2                               2                0                12
12                              1                1                8

      Product_Category_2  Product_Category_3  Purchase
2                9.842329                NaN        1422
12               9.842329                NaN        6073
```

```
[213]: preds = pd.DataFrame(y_pred_xgb).rename(columns={0: 'Predictions'}).
      ↪set_index(x_test.index)
preds.head(2)
```

```
[213]:      Predictions
24033    8125.541016
301904    7806.015137
```

```
[214]: output_df= output_df.merge(preds, how = 'left' ,left_index=True,
      ↪right_index=True)
output_df
```

```
[214]:      User_ID Product_ID Gender   Age  Occupation City_Category \
2    1000001  P00087842      F  0-17          10           A
12   1000005  P00031342      M 26-35          20           A
14   1000006  P00231342      F 51-55           9           A
22   1000008  P00213742      M 26-35          12           C
23   1000008  P00214442      M 26-35          12           C
...     ...         ...   ...   ...     ...
550044 1006004  P00370853      F 26-35          15           C
550047 1006009  P00372445      F 26-35          12           C
550048 1006010  P00371644      M 36-45           0           C
550051 1006013  P00375436      F 26-35          20           C
```

550064	1006035	P00375436	F	26-35	1	C
--------	---------	-----------	---	-------	---	---



	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
2	2	0	12	
12	1	1	8	
14	1	0	5	
22	4+	1	8	
23	4+	1	8	
...	
550044	2	0	19	
550047	3	0	20	
550048	1	0	20	
550051	3	0	20	
550064	3	0	20	

	Product_Category_2	Product_Category_3	Purchase	Predictions
2	9.842329	NaN	1422	1348.173950
12	9.842329	NaN	6073	7207.742676
14	8.000000	14.0	5378	6909.993164
22	9.842329	NaN	9743	7862.944824
23	9.842329	NaN	5982	7862.944824
...
550044	9.842329	NaN	62	245.385040
550047	9.842329	NaN	244	381.417999
550048	9.842329	NaN	591	420.123199
550051	9.842329	NaN	489	381.417999
550064	9.842329	NaN	371	381.417999

[110014 rows x 13 columns]

Let's visualize the predicted results.

```
[228]: sns.scatterplot(data=output_df.head(100), x='Purchase', y='Predictions')
plt.plot(y_pred_xgb, y_pred_xgb, 'r')
plt.show()
output_df[['Purchase', 'Predictions']].head(20).plot(kind='bar', figsize=(10, 8))
plt.show()
```



The Blue bar represents Actual purchase value and Orange bar represents Predicted purchase value. We can see that most of the Actual and predicted values are close. And model is Performing well and not overfitting.

[]:

**Predict on Unseen data.

[215]: `test_df.head(2)`

```
[215]:   Gender  City_Category  Stay_In_Current_City_Years  Marital_Status  \
0         1             2                             2                1
1         1             3                             0                0

   Product_Category_1  Product_Category_2  Age_Group_18-25  Age_Group_26-35  \
0                   1                  11.0              0                0
1                   3                   5.0              0                1

   Age_Group_36-45  Age_Group_46-50  Age_Group_51-55  Age_Group_55 & above
0                 0                 1                0                0
1                 0                 0                0                0
```

```
[217]: pred_xgb = xgb_reg.predict(test_df)
test_preds = pd.DataFrame(pred_xgb).rename(columns={0: 'Predicted Purchase'})
test_preds.head()
```

```
[217]:   Predicted Purchase
0      14242.577148
1      10929.032227
2       7042.043945
3       2967.486816
4       2478.158936
```

```
[218]: test_df.merge(test_preds, how = 'left', left_index = True, right_index = True)
```

```
[218]:   Gender  City_Category  Stay_In_Current_City_Years  Marital_Status  \
0         1             2                             2                1
1         1             3                             0                0
2         0             2                             4                1
3         0             2                             4                1
4         0             3                             1                0
...     ...           ...                           ...           ...
233594    0             2                             4                1
233595    0             2                             4                1
233596    0             2                             4                1
```

233597	0	3	4	0
233598	0	2	4	1

	Product_Category_1	Product_Category_2	Age_Group_18-25	\
0	1	11.000000	0	
1	3	5.000000	0	
2	5	14.000000	0	
3	4	9.000000	0	
4	4	5.000000	0	
...	
233594	8	9.849586	0	
233595	5	8.000000	0	
233596	1	5.000000	0	
233597	10	16.000000	0	
233598	4	5.000000	0	

	Age_Group_26-35	Age_Group_36-45	Age_Group_46-50	Age_Group_51-55	\
0	0	0	1	0	
1	1	0	0	0	
2	0	1	0	0	
3	0	1	0	0	
4	1	0	0	0	
...	
233594	1	0	0	0	
233595	1	0	0	0	
233596	1	0	0	0	
233597	0	0	1	0	
233598	0	0	1	0	

	Age_Group_55 & above	Predicted Purchase
0	0	14242.577148
1	0	10929.032227
2	0	7042.043945
3	0	2967.486816
4	0	2478.158936
...
233594	0	7344.817871
233595	0	6442.360352
233596	0	12347.626953
233597	0	19883.351562
233598	0	2308.947021

[233599 rows x 13 columns]